

# Threaded task assignment in CPLEX

Roberto Meroni\*      Hugo Sanz González†

November 2023

Consider the problem of distributing the workload required to perform a set of tasks  $T$  among a set of computers  $C$ . The workload of each task  $t \in T$  is divided into a set of threads  $H_t$ , and each  $h \in H_t$  requires  $r_h \geq 0$  compute units to complete. Each computer  $c \in C$  provides a set of cores  $K_c$ , all of which have  $r_c \geq 0$  compute units available. Consider a mixed integer linear program  $P_3$  where a thread is assigned to exactly one core, and all threads of a task are assigned to cores of a single computer. If the load of a computer  $c \in C$  refers to the sum of the loads of the cores in  $K_c$  divided by its total capacity  $|K_c| \cdot r_c$ , the objective in  $P_3$  is to minimize the load of the highest-loaded computer.

## 1. Implement the $P_3$ model in OPL and solve it using CPLEX.

Let  $H$  be the set of all threads among all tasks in  $T$ , and let  $K$  be the set of all cores among all computers in  $C$ . For each thread  $h \in H$  and each core  $k \in K$ , define a binary variable  $x_{hk}$  denoting whether  $h$  is assigned to  $k$ . Also let  $x_{tc}$  be a binary variable denoting whether the threads of a task  $t \in T$  are assigned to the cores of a computer  $c \in C$ . The corresponding variable definitions in the OPL language look as follows:

```
range H = 1..nThreads;  
range C = 1..nCPUs;  
float rh[h in H] = ...;  
float rc[c in C] = ...;
```

All threads must be assigned to exactly one core, hence the first set of constraints is  $\sum_{k \in K} x_{hk} = 1$  for all  $h \in H$ . In OPL this can be expressed as

```
forall(h in H)  
    sum(k in K) x_hk[h, k] == 1;
```

where the sets  $H$  and  $K$  are declared as

```
range H = 1..nThreads;  
range K = 1..nCores;
```

In addition, all the threads belonging to a task must be assigned to cores belonging to the same computer. If the threads of a task  $t \in T$  are assigned to some subset of cores of a computer  $c \in C$ , then  $|H_t|x_{tc}$  is equal to the number of threads of  $t$ ; otherwise  $x_{tc}$  is zero and so this last quantity is also

---

\*roberto.meroni@estudiantat.upc.edu

†hugo.sanz@estudiantat.upc.edu

nil. Since  $\sum_{h \in H_t, k \in K_c} x_{hk}$  counts the number of threads belonging to task  $t$  that are assigned to cores in  $c$ , the constraints

$$\sum_{h \in H_t} \sum_{k \in K_c} x_{hk} = |H_t| x_{tc}, \quad \forall t \in T, c \in C$$

ensure that the threads of a task are all assigned to cores of a single computer. In OPL,

```
forall(c in C, t in T)
    sum(h in H, k in K) x_hk[h, k] * CK[c, k] * TH[t, h]
    == x_tc[t, c] * sum(h in H) TH[t, h];
```

The entry  $i, j$  of matrix CK is a boolean value describing whether a core  $k_j \in K$  is part of a computer  $c_i \in C$ . Similarly entry  $i, j$  of matrix TH defines whether a thread  $h_j \in H$  forms part of task  $t_i \in T$ . We declare them in OPL as follows:

```
float CK[c in C, k in K] = ...;
float TH[t in T, h in H] = ...;
```

Another set of constraints ensures that the capacity of all cores  $K_c$  in a computer  $c \in C$  is not exceeded. Note that  $r_h x_{hk}$  is the load assigned to a core  $k \in K_c$  corresponding to a thread  $h \in H$ , so

$$\sum_{h \in H} r_h x_{hk} \leq r_c, \quad \forall c \in C, k \in K_c.$$

In OPL this is

```
forall (c in C, k in K)
    sum(h in H) rh[h] * x_hk[h, k] * CK[c, k] <= rc[c];
```

The total number of compute units handled by all cores of a computer  $c \in C$  is  $\sum_{h \in H} \sum_{k \in K_c} r_h x_{hk}$ , and the capacity of  $c$  is its number of cores  $|K_c|$  multiplied by the capacity  $r_c$  of each one of them. Introduce a real decision variable  $z$  representing the load of the highest-loaded computer. This is precisely the nonnegative quantity to be minimized, and it is bounded below by the load of each computer. That is,

$$z \geq \frac{1}{|K_c| r_c} \sum_{h \in H} \sum_{k \in K_c} r_h x_{hk}, \quad \forall c \in C. \quad (1)$$

In OPL, this set of constraints may be written as

```
forall(c in C)
    z >= (sum(h in H, k in K) rh[h] * x_hk[h, k] * CK[c, k])
    / (rc[c] * sum(k in K) CK[c, k]);
```

At this point problem  $P_3$  has been fully specified, and we may solve it for the particular instance specified in the provided data file. We have computers  $c_1, c_2, c_3$  with cores  $k_{i1}, k_{i2}, k_{i3}$  for each  $c_i \in C$ , and tasks  $t_1, \dots, t_4$  with threads  $h_{i1}, h_{i2}$  for each  $t_i \in T$ . The CPLEX solver finds the assignment

$$\begin{array}{cccc} h_{11} \rightarrow k_{12} & h_{21} \rightarrow k_{11} & h_{31} \rightarrow k_{31} & h_{41} \rightarrow k_{32} \\ h_{12} \rightarrow k_{22} & h_{22} \rightarrow k_{21} & h_{32} \rightarrow k_{11} & h_{42} \rightarrow k_{11} \end{array}$$

where  $h_{ij} \rightarrow k_{pq}$  means thread  $h_{ij}$  of task  $t_i \in T$  is assigned to core  $k_{pq}$  of computer  $c_p \in C$ .

2. *Generate instances of increasing size and use the  $P_3$  model to solve them.*

We modify the provided InstanceGeneratorP3/config.dat file in order to produce five instances having the following characteristics:

#	Tasks	Threads		Computers	Cores	
		Per task	Total		Per computer	Total
1	10	[10, 15]	128	5	[100, 150]	585
2	10	[10, 20]	152	10	[400, 700]	5383
3	10	[5, 25]	126	5	[500, 1000]	3789
4	10	[5, 25]	115	5	[750, 1500]	4801
5	12	[5, 10]	87	5	[50, 100]	362

The number of threads per task and the number of cores per computer are random integers drawn from the corresponding intervals given in the table. The capacity per core is 5000 and the compute units required by each thread follow a uniform distribution in  $[500, 800]$ . Details regarding the problem size and the solutions found by CPLEX appear in Exercise 4.

3. *Modify the  $P_3$  model to maximize the number of computers with all their cores empty.*

Introduce a binary decision variable  $y_c$  for each computer  $c \in C$  that denotes whether no thread has been assigned to any of the cores in  $K_c$ . If  $y_c = 0$ , then there exists a task  $t \in T$  such that  $x_{tc} = 1$ . In other words,  $y_c = 0$  implies that  $\sum_{t \in T} x_{tc} \geq 1$ , or equivalently  $1 - \sum_{t \in T} x_{tc} \leq 0$ . Since the number 1 is an upper bound on the left-hand side of this equation, the set of constraints that express these implications is

$$1 - \sum_{t \in T} x_{tc} \leq y_c, \quad \forall c \in C.$$

In OPL this is expressed as follows:

```
forall(c in C)
  1 - sum(t in T) x_tc[t, c] <= y_c[c];
```

where the  $y_c$ 's are declared as

```
dvar boolean y_c[c in C];
```

On the other hand, if  $y_c = 1$  then  $x_{tc} = 0$  for all  $t \in T$ . Since  $|T|$  is an upper bound on  $\sum_{t \in T} x_{tc}$ , we want

$$\sum_{t \in T} x_{tc} \leq |T|(1 - y_c), \quad \forall c \in C.$$

Analogously we may write this in OPL as follows:

```
forall(c in C)
  sum(t in T) x_tc[t, c] <= nTasks * (1 - y_c[c]);
```

#	Model	Variables	Constraints	Iterations	Runtime (s)
1	$P_3$	74 931	3108	429 390	28
	$P'_3$	74 935	3113	462	3
2	$P_3$	818 317	54 092	66 339	152
	$P'_3$	818 326	54 102	663	47
3	$P_3$	477 465	19 126	881 293	828
	$P'_3$	477 469	19 131	424	23
4	$P_3$	552 166	24 175	816 669	711
	$P'_3$	552 170	24 180	419	26
5	$P_3$	31 555	1962	1 203 320	44
	$P'_3$	31 559	1967	335	1

Table 1: The statistics of the instances generated in exercise 2 under the  $P_3$  and  $P'_3$  models, as well as the solution time achieved by CPLEX.

Our goal is to maximize the number of computers with no work assigned, so we remove the constraint in Eq. (1) and replace the objective in  $P_3$  by  $\max \sum_{c \in C} y_c$  to obtain the new model  $P'_3$ .

4. Compare models  $P_3$  and  $P'_3$  in terms of number of variables, number of constraints and execution time for the generated instances.

The problem size and results found by CPLEX after solving the instances generated in Exercise 2 under the two models appear in Table 1; all problems are feasible. In  $P_3$ , the appearance of a continuous decision variable makes the problem a mixed-integer linear program and this increases the time needed by CPLEX to find an optimal solution. In contrast,  $P'_3$  is a binary linear program which is more amenable to branch & cut techniques. For example, for instances 1, 3 and 5 the time decreased respectively by a factor of 9.33, 36 and 44.

It is interesting to compare the solver's behavior for instances 1 and 5; even if the former has roughly twice as many variables and constraints, the latter takes  $\approx 1.57$  times longer to solve under the  $P_3$  model. This suggests that the total number of variables and constraints are not the only factor affecting the solving time; the shape of the problem also has a profound effect. (Here instance 1 has 2 less tasks to assign than instance 5.)

We could not produce instances that lead to solving times greater than 30 minutes due to memory limitations. Still, the results indicate that minimizing the load of the highest-loaded computer is a much harder problem to solve than maximizing the number of computers with no assigned tasks.